

An Introduction to Real-time Digital Signal Processing through Parallelism

Sherwyn Sen
CS 159





DISCLAIMER


The code I will show in the presentation is not the actual code. It is simplified for the sake of explanation.

REAL CODE (feel free to look at this during the presentation)

https://github.com/ika-musuko/parallelizing-dsp/tree/master/wave_analyzer


What is Digital Signal Processing?

- Take a signal and modify it
- Signals
 - Motion
 - Sound
 - Image Data





What is Digital Signal Processing?

- Take a signal and modify it
- Signals
 - Motion
 - **Sound**
 - Easiest to explain
 - Image Data





How Signals Get Processed By Computers




How Signals Get Processed By Computers



How Signals Get Processed By Computers



How Signals Get Processed By Computers




How are we going to Parallelize DSP?

- Let's make a spectrum analyzer! (in software)



What is a Spectrum Analyzer?

- Measures magnitude of an input signal based on its frequencies



How does a Spectrum Analyzer work?

- Fourier Transform



How does a Spectrum Analyzer work?

- Fourier Transform

$$X_k = \frac{1}{N} \sum_{n=0}^{N-1} x_n e^{i2\pi k \frac{n}{N}}$$



How does a Spectrum Analyzer work?

- Fourier Transform

$$X_k = \frac{1}{N} \sum_{n=0}^{N-1} x_n e^{i2\pi k \frac{n}{N}}$$

To find the energy at a particular frequency, spin your signal around a circle at that frequency, and average a bunch of points along that path.




How does a Spectrum Analyzer work?

- Fourier Transform

$$X_k = \frac{1}{N} \sum_{n=0}^{N-1} x_n e^{i2\pi k \frac{n}{N}}$$

To find the energy at a particular frequency, spin your signal around a circle at that frequency, and average a bunch of points along that path.

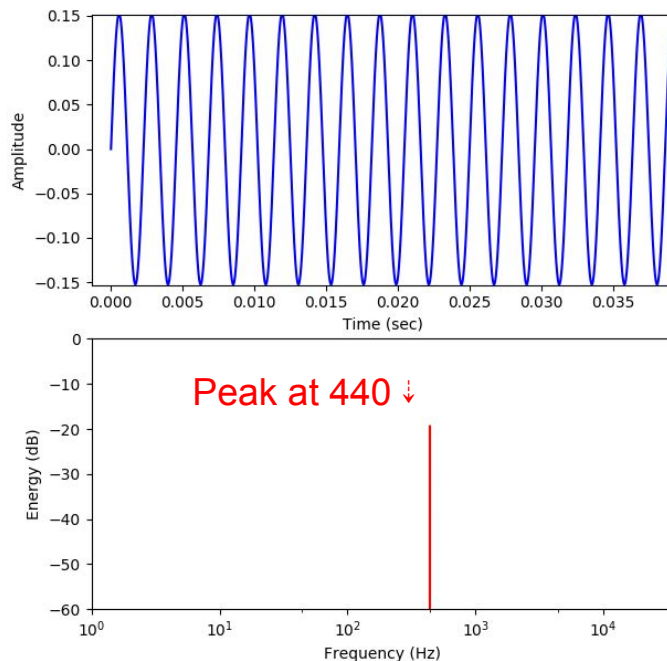
For a more in-depth explanation of the Fourier Transform, consult But what is the Fourier Transform? A visual introduction by 3blue1brown



Basic Spectrum Analysis (not real time)


440 Hz Sine Wave

“Result of the Fourier Transform”



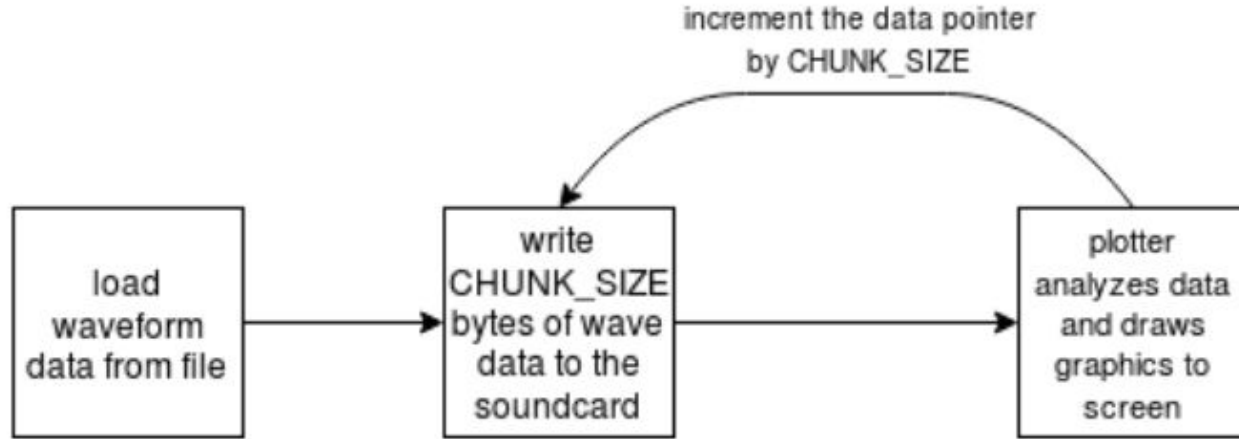
`fft_plotter.py`

Basic Spectrum Analysis (not real time)



fft_plotter.py

Simple approach to attempt real-time analysis




Simple approach to attempt real-time analysis

```
def play(wf: wave_file):  
    # read the first 1024 bytes of the file  
    stream = Stream()  
    plotter = Plotter()  
    data = wf.readframes(CHUNK_SIZE)  
  
    # while aren't at the end of the file  
    while len(data) > 0:  
  
        # write the data to the soundcard  
        stream.write(data)  
  
        # analyze the buffer and draw the  
        # graphics to the screen  
        # this function uses the FFT described  
        # earlier to do so  
        plotter.plot(data)  
  
        # read the next 1024 bytes of wave  
        # data  
        data = wf.readframes(CHUNK_SIZE)
```

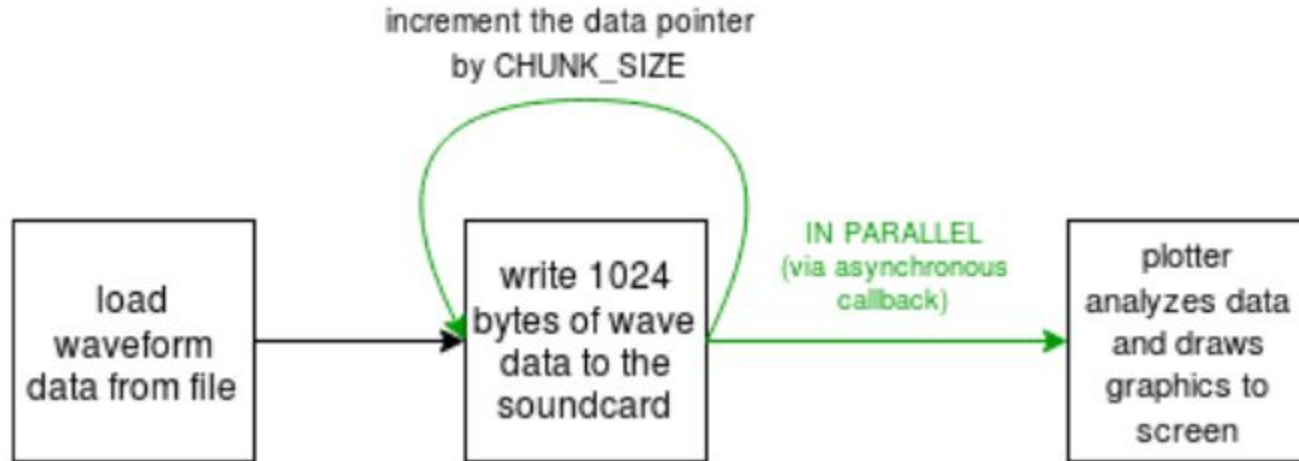


okay! let's try it...

What happened???



Basic Parallel Analysis - Strategy



Basic Parallel Analysis - Strategy

Splitting the player and plotter into processes

```
plot_proc = Process(target=start_plotter)
plot_proc.start()
play()
```

Basic Parallel Analysis Strategy - Problems


```
def play(wf: wave_file):  
    # read the first 1024 bytes of the file  
    stream = Stream()  
    plotter = Plotter()  
    data = wf.readframes(CHUNK_SIZE)  
  
    # while aren't at the end of the file  
    while len(data) > 0:  
        # write the data to the soundcard  
        stream.write(data)  
  
        # analyze the buffer and draw the  
        # graphics to the screen  
        # this function uses the FFT described  
        # earlier to do so  
        plotter.plot(data)  
  
        # read the next 1024 bytes of wave  
        # data  
        data = wf.readframes(CHUNK_SIZE)
```

!!! These functions are
blocking I/O !!! Can't use
with parallelism !!!!

bad_analyzer.py

The Solution? Asynchronous Callback Functions

- Callback functions
 - Functions passed to another function or object
 - Called when some sort of event happens
 - Asynchronous callbacks done in parallel using threads/processes
 - Well known examples: Javascript setTimeout/setInterval



The Solution? Asynchronous Callback Functions

- “Stream” example

```
# stream: a nonblocking I/O "stream"
class Stream:
    def __init__(callback_function, args):
        self.callback_function = callback_function
        self.args = args

    def start_stream(self):
        p = Process(
            target = self.callback_func, args = self.args)
        p.start()
```

for audio stream example, see PortAudio's portaudio_startstream.c

The Solution? Asynchronous Callback Functions

audio playback

```
def get_wave_data(wf: wave_file):  
    # read the next set of wave data bytes  
    data = wf.readframes(CHUNK_SIZE)  
  
    # put the wave data onto the shared queue  
    (for the animation)  
    SHARED_DATA_QUEUE.put(data)  
  
    # the data the Stream object should deal with  
    return data  
  
def play(wf: wave_file):  
    stream = Stream(  
        callback_function=get_wave_data, args=wf)  
    stream.start_stream()  
    while stream.is_active():  
        time.sleep(0.1)
```

better_analyzer.py

The Solution? Asynchronous Callback Functions

plotting on the graph


```
def plot_callback():  
    # if there's nothing on the data queue,  
    just make the data all 0  
    if self.data_queue.empty():  
        data = [0]*CHUNK_SIZE  
  
    # otherwise get the data from the data_queue  
    else:  
        data = SHARED_DATA_QUEUE.get()  
  
    # the data the Plotter object should deal with  
    return data  
  
# initialize the plotter animation callback  
def start_plotter():  
    plotter = Plotter(callback_function=plot_callback)  
    plotter.start_plotting()
```

better_analyzer.py

Shared Data Queue - Audio Playback Side

audio playback

```
def get_wave_data(wf: wave_file):  
    # read the next set of wave data bytes  
    data = wf.readframes(CHUNK_SIZE)  
  
    # put the wave data onto the shared queue  
    (for the animation)  
    SHARED_DATA_QUEUE.put(data)  
  
    # the data the Stream object should deal with  
    return data  
  
def play(wf: wave_file):  
    stream = Stream(  
        callback_function=get_wave_data, args=wf)  
    stream.start_stream()  
    while stream.is_active():  
        time.sleep(0.1)
```




PUT audio data
chunks onto the queue
before playing them

better_analyzer.py

Shared Data Queue - Graph Plotting Side

plotting on the graph

```
def plot_callback():  
    # if there's nothing on the data queue,  
    # just make the data all 0  
    if self.data_queue.empty():  
        data = [0]*CHUNK_SIZE  
  
    # otherwise get the data from the data_queue  
    else:  
        data = SHARED_DATA_QUEUE.get()  
  
    # the data the Plotter object should deal with  
    return data  
  
# initialize the plotter animation callback  
def start_plotter():  
    plotter = Plotter(callback_function=plot_callback)  
    plotter.start_plotting()
```



GET audio data
chunks from the queue
and graph them

QUESTION



Why not just use
a shared data
variable between
the two
processes?

ROUND 2



Almost there...What's the problem?

- The **chunk producing audio player** is much faster than the **chunk consuming graph plotter**
- Cause of the LAG!!



Solution: Queue Eating

```
data = SHARED_DATA_QUEUE.get()
while not SHARED_DATA_QUEUE.empty():
    data = SHARED_DATA_QUEUE.get()
```


- Plotter pops off all of the items in the queue (“eats” the items)
 - Relatively fast operation compared to processing the chunk
- Only processes the last data chunk on the queue
- *The screen needs to show “now”, not 20 chunks ago!*

The background is a solid pink color. In the top right corner, there is a geometric pattern consisting of several squares and triangles in different shades of pink, creating a stepped or architectural effect.

WILL THIS WORK?

Conclusion

- Took a basic digital signal processing routine
- Identified the bottleneck
- Partitioned the tasks
- Parallelized the partitioned tasks



What can be done from here?

- Use multiple wave analyzers, each with their own inputs to create a simple mixer
 - The “master” channel could have its own analyzer as well
- Create an interface for applying effects to the audio in real time
 - Filters
 - Delays
 - Reverb
- Possibilities are endless

